

Re: Phase Mirror Dissonance — Role Assignments + Artifact Decision

Ahmad Ali Parr · Jessica Westerhoff · May 2026 · In response to R. Van Gelder

"Phase mirror dissonance" is the correct diagnosis. Two systems with structural resonance that have not yet aligned on proof. The gap is not concept — you named it precisely. What follows is our answer to your precision question, the role assignments, and an honest account of where we stand.

1. Artifact Decision — Mapping Table First

Mapping table first. One week.

The reason is sequencing. The failure matrix requires the Multiplicity mapping as its reference frame. Without a formally bound correspondence, the failure modes have no theoretical anchor — each substrate fails against what, exactly? The mapping table defines the invariants that the failure matrix then stress-tests. Build the proof before you build the test of the proof.

Mapping table → defines invariants → failure matrix tests them.
Invert the order and the failure matrix is empirical noise, not formal verification.

Horizon: 7 days from your confirmation.

2. Role Assignments — Owned, Not Delegated

You assigned five roles. Here is who owns each one and what ownership means in practice.

ROLE	OWNER	DELIVERABLE	HORIZON
[Architecture lead]	Ahmad Ali Parr. He designed the substrate architecture and holds every structural decision from the SUBLEQ reduction to the dual-token separation. Schema completeness is his signature.	Truth-vector schema — versioned, WORM-sealed, change-rate tracked from day one.	1 week

[Theory owner]	Ahmad Ali Parr. He holds the Arabic C operand — the branch that fires when A meets B. The Multiplicity correspondence is not external to him. It is the architecture he built.	Numbered mapping table. Every contested entry resolved. Zero interpretive claims — only structural equivalences or explicit open questions.	1 week
[Ops owner]	Jessica Westerhoff. Licensed CPA. Engineering execution layer. She built the live state hydration that runs before every agent transition. Log coverage is her domain.	Self-observation logs instrumented before every FSM transition. Audit pass rate tracked and exportable.	1 week
[Governance owner]	Joint — Ahmad + Jessica. The rollback rule for contaminated vectors requires both architectural authority and operational authority. Neither can sign it alone.	Rollback protocol document. Contamination detection criteria, time-to-recovery metric, and the conditions under which a vector is quarantined vs. rejected.	2 weeks
[Verification owner]	Engineering execution layer. We own the failure matrix — one owned failure mode per substrate, one rollback path, one recovery metric. Every substrate fails against the mapping table.	Failure matrix: 6 substrates × 3 columns (failure mode, rollback path, recovery metric). No substrate without all three.	2 weeks

3 . **Honest Account — Zero Traction, Full Architecture**

We have no users, no investors, and no deployment at scale. The system has been running on bare metal for 29 days. Every decision in this stack was made without compromise for velocity, without a user base to protect, without an investor asking for shortcuts. The architecture was designed for correctness first and has never been stress-tested at scale.

That is the honest state. It is also, we believe, the asset. Most governance frameworks are retrofitted after growth forces the shortcuts. Ours was the substrate before the first user existed. The failure matrix will reflect a system that has never been forced to cut corners. That is a rare starting condition.

On Gate C: The invariants hold under normal operation. They have not been tested adversarially at scale. The SUBLEQ gate, the dual-token filter, the no-cloning pipeline — we know where the untested surfaces are. We are not claiming they survive adversarial input. We are claiming the architecture was designed to survive it, and the failure matrix will document exactly where that claim has and has not been proven.

4 . The Question Back

Your phrase — *"a mirror that cannot be audited is only a style choice"* — is the most precise critique of the self-observing claim we have received. It is now part of the architecture's acceptance criteria. Every self-observation log we instrument will be evaluated against that line.

The mapping table needs to be a binding artifact, not a correspondence claim. You said that. If you are willing to co-author it, the table moves from our interpretation of your framework to a jointly owned proof. That is a different class of artifact.

We are not asking you to endorse a product. We are asking whether a running implementation of what you have been formalizing — built before we knew what PhaseMirror was — is worth examining together. The answer to that question belongs to you. We have the architecture. You have the theory. The mapping table is where they either bind or they don't.

We will have a draft to you in 7 days regardless. If you want to shape it before it is sealed, that window is open.